

CS 5350/6350: Machine Learning Fall 2016


Homework 1

Gopal Menon


September 12, 2016

1 Decision trees (35 points)


1. (a) $(x_1 \vee x_2) \wedge x_3$



- (b) $(x_1 \wedge x_2) \text{ xor } (\neg x_1 \vee x_3)$



- (c) The 2-of-3 function defined as follows: at least 2 of $\{x_1, x_2, x_3\}$ should be true for the output to be true.



2. (a) [2 points] How many possible functions are there to map these four features to a Boolean decision?

The number of possible rows in a truth table with all possible values of the four features will be

$$\begin{aligned} & \text{Number of types of Berries} \times \text{Number of types of Balls} \times \text{Number of colors} \times \\ & \text{Number of Pokémon types} \\ & = 2 \times 3 \times 3 \times 4 = 72 \end{aligned}$$

Each row in this truth table can have a value of *Yes* or *No* as a label for whether the Pokémon can be caught or not. The *Yes* and *No* can be represented as bits 1 and 0. Each possible function will map to one combination of these labels and will be of the form of a 72-bit binary number. So the total possible number of functions will be 2^{72} , which is the count of all the possible 72-bit binary numbers.

- (b) [2 points] What is the entropy of the labels in this data? (When calculating entropy, The base of the logarithm should be 2.)

The entropy of a collection S where the target label can take on c different values is defined as [1]

$$Entropy(S) = \sum_{i=1}^c -p_i \log_2 p_i$$

where p_i is the proportion of S belonging to label i .

$$\begin{aligned} Entropy(PokémonData) &= -\frac{8}{16} \log_2 \frac{8}{16} - \frac{8}{16} \log_2 \frac{8}{16} \\ &= -\frac{1}{2} \log_2 \frac{1}{2} - \frac{1}{2} \log_2 \frac{1}{2} \\ &= -\log_2 \frac{1}{2} \\ &= 1 \end{aligned}$$

- (c) [8 points] Calculate information gain for four features respectively. Keep 3 significant digits.

$$Values(Berry) = Yes, No$$

$$S = [8+, 8-]$$

$$S_{Yes} = [6+, 1-]$$

$$S_{No} = [2+, 7-]$$

$$\begin{aligned} Gain(S, Berry) &= Entropy(S) - \frac{7}{16}Entropy(S_{Yes}) - \frac{9}{16}Entropy(S_{No}) \\ &= 1 - \frac{7}{16} \times 0.5917 - \frac{9}{16} \times 0.7642 \\ &= 0.311 \end{aligned}$$

$$Values(Ball) = Poké, Great, Ultra$$

$$S_{Poké} = [1+, 5-]$$

$$S_{Great} = [4+, 3-]$$

$$S_{Ultra} = [3+, 0-]$$

$$\begin{aligned} Gain(S, Ball) &= Entropy(S) - \frac{6}{16}Entropy(S_{Poké}) - \frac{7}{16}Entropy(S_{Great}) - \frac{3}{16}Entropy(S_{Ultra}) \\ &= 1 - \frac{6}{16} \times 0.65 - \frac{7}{16} \times 0.9852 - \frac{3}{16} \times 0 \\ &= 0.3252 \end{aligned}$$

$$Values(Color) = Green, Yellow, Red$$

$$S_{Green} = [2+, 1-]$$

$$S_{Yellow} = [3+, 4-]$$

$$S_{Red} = [3+, 3-]$$

$$\begin{aligned} Gain(S, Color) &= Entropy(S) - \frac{3}{16}Entropy(S_{Green}) - \frac{7}{16}Entropy(S_{Yellow}) \\ &\quad - \frac{6}{16}Entropy(S_{Red}) \\ &= 1 - \frac{3}{16} \times 0.9183 - \frac{7}{16} \times 0.9852 - \frac{6}{16} \times 1 \\ &= 0.0218 \end{aligned}$$

$$Values(Type) = Normal, Water, Flying, Psychic$$

$$S_{Normal} = [3+, 3-]$$

$$S_{Water} = [2+, 2-]$$

$$S_{Flying} = [3+, 1-]$$

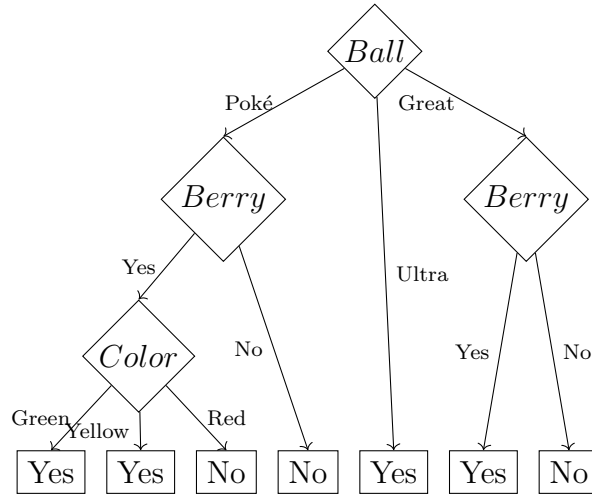
$$S_{Psychic} = [0+, 2-]$$

$$\begin{aligned}
Gain(S, Type) &= Entropy(S) - \frac{6}{16} Entropy(S_{Normal}) - \frac{4}{16} Entropy(S_{Water}) \\
&\quad - \frac{4}{16} Entropy(S_{Flying}) - \frac{2}{16} Entropy(S_{Psychic}) \\
&= 1 - \frac{6}{16} \times 1 - \frac{4}{16} \times 1 - \frac{4}{16} \times 0.8113 - \frac{2}{16} \times 0 \\
&= 0.172
\end{aligned}$$

- (d) [3 points] According to your results, using ID3 algorithm which attribute should be root for the decision tree?

The root should be feature *Ball*, since it has the largest entropy gain.

- (e) [4 points] Construct a decision with the root you selected in the previous question. You do not have to use the ID3 algorithm here, you can show any tree with the chosen root.



- (f) [2 points] Using your decision tree to predict label in the test set in the table below, what is your label for the each example? What is your accuracy?

Only one out of three was classified correctly. So accuracy is low.

Berry	Ball	Color	Type	Caught	Prediction
Yes	Great	Yellow	Psychic	Yes	Yes
Yes	Poké	Green	Flying	No	Yes
No	Ultra	Red	Water	No	Yes

- (g) [1 points] Do you think it is a good idea to use decision tree in this Pokémon Go problem?

It is not a good idea to use a decision tree for this particular Pokémon Go problem. The training data looks to be not enough for a learning algorithm and the test data seems to be adversarial. A decision tree may perform better with a larger test data set.

3. Recall that in the ID3 algorithm, we want to identify the best attribute that splits the examples that are relatively pure in one label. Apart from entropy, which you used in

the previous question, there are other methods to measure impurity. One such impurity measure is the Gini measure, that is used in the CART family of algorithms. If there are k possible outcomes $1, \dots, i, \dots, k$, each with a probability $p_1, \dots, p_i, \dots, p_k$ of occurring, the Gini measure is defined as:

$$Gini(p_1, \dots, p_k) = 1 - \sum_{i=1}^k p_i^2$$

The Gini measure can be used to replace entropy in the definition of information gain to pick the best attribute.

- (a) [4 points] Using the Gini measure, calculate the information gain for the four features respectively. Use 3 significant digits.

$$\begin{aligned} Gini(PokémonData) &= 1 - \left(\frac{8}{16}\right)^2 - \left(\frac{8}{16}\right)^2 \\ &= 1 - 0.25 - 0.25 \\ &= 0.5 \end{aligned}$$

$$\begin{aligned} Gain(S, Berry) &= Gini(S) - \frac{7}{16}Gini(S_{Yes}) - \frac{9}{16}Gini(S_{No}) \\ &= 0.5 - \frac{7}{16} \left(1 - \left(\frac{6}{7}\right)^2 - \left(\frac{1}{7}\right)^2\right) - \frac{9}{16} \left(1 - \left(\frac{2}{9}\right)^2 - \left(\frac{7}{9}\right)^2\right) \\ &= 0.5 - \frac{7}{16} \times \frac{12}{49} - \frac{9}{16} \times \frac{28}{81} \\ &= 0.5 - 0.1071 - 0.1944 \\ &= 0.198 \end{aligned}$$

$$\begin{aligned} Gain(S, Ball) &= Gini(S) - \frac{6}{16}Gini(S_{Poké}) - \frac{7}{16}Gini(S_{Great}) - \frac{3}{16}Gini(S_{Ultra}) \\ &= 0.5 - \frac{6}{16} \left(1 - \left(\frac{1}{6}\right)^2 - \left(\frac{5}{6}\right)^2\right) - \frac{7}{16} \left(1 - \left(\frac{4}{7}\right)^2 - \left(\frac{3}{7}\right)^2\right) \\ &\quad - \frac{3}{16} \left(1 - \left(\frac{3}{3}\right)^2 - \left(\frac{0}{3}\right)^2\right) \\ &= 0.5 - \frac{6}{16} \times \frac{10}{36} - \frac{7}{16} \times \frac{24}{49} - \frac{3}{16} \times 0 \\ &= 0.5 - 0.1042 - 0.2143 \\ &= 0.181 \end{aligned}$$

$$\begin{aligned}
Gain(S, Color) &= Gini(S) - \frac{3}{16}Gini(S_{Green}) - \frac{7}{16}Gini(S_{Yellow}) - \frac{6}{16}Gini(S_{Red}) \\
&= 0.5 - \frac{3}{16} \left(1 - \left(\frac{2}{3} \right)^2 - \left(\frac{1}{3} \right)^2 \right) - \frac{7}{16} \left(1 - \left(\frac{3}{7} \right)^2 - \left(\frac{4}{7} \right)^2 \right) \\
&\quad - \frac{6}{16} \left(1 - \left(\frac{3}{6} \right)^2 - \left(\frac{3}{6} \right)^2 \right) \\
&= 0.5 - \frac{3}{16} \times \frac{4}{9} - \frac{7}{16} \times \frac{24}{49} - \frac{6}{16} \times \frac{1}{2} \\
&= 0.5 - 0.0833 - 0.2143 - 0.1875 \\
&= 0.015
\end{aligned}$$

$$\begin{aligned}
Gain(S, Type) &= Gini(S) - \frac{6}{16}Gini(S_{Normal}) - \frac{4}{16}Gini(S_{Water}) - \frac{4}{16}Gini(S_{Flying}) \\
&\quad - \frac{2}{16}Gini(S_{Psychic}) \\
&= 0.5 - \frac{6}{16} \left(1 - \left(\frac{3}{6} \right)^2 - \left(\frac{3}{6} \right)^2 \right) - \frac{4}{16} \left(1 - \left(\frac{2}{4} \right)^2 - \left(\frac{2}{4} \right)^2 \right) \\
&\quad - \frac{4}{16} \left(1 - \left(\frac{3}{4} \right)^2 - \left(\frac{1}{4} \right)^2 \right) - \frac{2}{16} \left(1 - \left(\frac{0}{2} \right)^2 - \left(\frac{2}{2} \right)^2 \right) \\
&= 0.5 - \frac{6}{16} \times \frac{1}{2} - \frac{4}{16} \times \frac{1}{2} - \frac{4}{16} \times \frac{6}{16} - \frac{2}{16} \times 0 \\
&= 0.5 - 0.1875 - 0.125 - 0.09375 \\
&= 0.094
\end{aligned}$$

- (b) [3 points] According to your results in the last question, which attribute should be the root for the decision tree? Do these two measures (entropy and Gini) lead to the same tree?

As per the results in the last question based on Gini calculation, the attribute *Berry* must be the root for the decision tree as it has the maximum gain. In the case of Entropy calculation, the attribute *Ball* was the root of the decision tree. So the two measures will lead to different trees.

2 Linear Classifiers (15 points)

In the questions in this section, we have four features x_1 , x_2 , x_3 and x_4 and the label is represented by o .

1. [3 points] Write a linear classifier that correctly classifies the given dataset. You don't need to run any learning algorithm here. Try to find the weights and the bias of the classifier using the definition of linear separators.

x_1	x_2	x_3	x_4	o
0	0	0	1	1
0	0	1	1	1
0	0	0	0	-1

I obtained the weights and bias by trial and error.

$$w = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 1 \end{bmatrix}$$

$$b = -1$$

2. [5 points] Suppose the dataset below is an extension of the above dataset. Check if your classifier from the previous question correctly classifies the dataset. Report its accuracy.

x_1	x_2	x_3	x_4	o
1	0	1	1	1
0	1	0	1	1
1	0	1	0	1
1	1	0	0	1
1	1	1	1	1
1	1	1	0	1
0	0	1	0	-1

Using the above classifier, the output is

<i>Classifier Output</i>	<i>Expected Output</i>
1	1
1	1
1	1
-1	1
1	1
1	1
1	-1

The classifier correctly classified 5 out of 7 outputs, giving an accuracy of 71.43%.

3. [7 points] Given the remaining missing data points of the above dataset in the table below, find a linear classifier that correctly classifies the whole dataset (all three tables together).

x_1	x_2	x_3	x_4	o
0	1	0	0	-1
0	1	1	0	-1
0	1	1	1	1
1	0	0	0	1
1	0	0	1	1
1	1	0	1	1

Upon writing down the entire dataset, it becomes evident that the output depends only on the value of x_1 . When $x_1 = 0$, the output is -1 and when $x_1 = 1$, the output is 1 . When that is the case, the weight and bias will be

$$w = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$b = -1$$

3 Experiments

Setting A [25 points]

1. [10 points] **Implementation**

- (a) [4 points] I used the Pokémon training and test data to verify that the decision tree had been coded properly. The decision tree was coded based on the algorithm in Tom Mitchell's book [1]. A random number generator was used when I needed to break a tie in the selection of the most common value of an attribute. To check the accuracy of the classifier I used the statistical metrics for Precision, Recall, Accuracy and F1 Score. The Java Collection library classes were used as they made programming a little easier. The Tree was implemented using a class hierarchy of generic node as the parent class of a leaf node and an internal node. The child nodes for an internal node were represented using a collection of nodes.
- (b) [2 points] There was no error when testing on the training data.
- (c) [5 points] There was no error when testing on the testing data.
- (d) [1 points] The maximum tree depth was 3.

2. [15 points] **Limiting Depth**

- (a) [10 points]

<i>Depth</i>	<i>Average Accuracy</i>	<i>StandardDeviation</i>
1	0.983	0.037
2	0.997	0.007
3	1.0	0.0
4	1.0	0.0
5	1.0	0.0
10	1.0	0.0
15	1.0	0.0
20	1.0	0.0

- (b) [5 points] A maximum depth of 3 was specified for the decision tree since it was the smallest tree that resulted in an average accuracy of 1.0 with 6-fold cross validation. With training and test data and a depth of 3, the accuracy obtained was 1.0.

Setting B [25 points]

1. [10 points] Experiments

For this problem, you will be using the data found in the `SettingB` folder. This folder contains the two files, `SettingB/training.data` and `SettingB/test.data`. In this setting you will be training your algorithm on the training file (`SettingB/training.data`). Remember that you should not look at or use your testing file until your algorithm is complete. You are not limiting the depth of your tree in this section.

- (a) [2 points] Report the error of your decision tree on the `SettingB/training.data` file.
- (b) [2 points] Report the error of your decision tree on the `SettingB/test.data` file.
- (c) [2 points] Report the error of your decision tree on the `SettingA/training.data` file.
- (d) [2 points] Report the error of your decision tree on the `SettingA/test.data` file.
- (e) [1 points] Report the maximum depth of your decision tree.

2. [15 points] Limiting Depth

In this section you will be using 6-fold cross-validation in order to limit the depth of your decision tree, effectively pruning the tree to avoid overfitting. You will be using the 6 cross-validation files for this section, titled `SettingB/CVSpplits/training_0X.data` where `X` is a number between 0 and 5 (inclusive).

- (a) [10 points] Run 6-fold cross-validation using the specified files. Experiment with depths in the set $\{1, 2, 3, 4, 5, 10, 15, 20\}$, reporting the cross-validation accuracy and standard deviation for each depth. Explicitly specify which depth should be chosen as the best, and explain why.
- (b) [5 points] Using the depth with the greatest cross-validation accuracy from your experiments: train your decision tree on the `SettingB/training.data` file. Report the accuracy of your decision tree on the `SettingB/test.data` file.

Setting C (CS 6350 Students) [20 points]

In this setting, you are investigating what effect missing features have on a decision tree, and exploring which approach is most effective in dealing with missing features. More specifically, you will be trying:

- **Method 1:** Setting the missing feature as the majority feature value.
- **Method 2:** Setting the missing feature as the majority value of that label.
- **Method 3:** Treating the missing feature as a *special* feature.

The missing feature is represented by a ? character. In order to determine which method is the best, you will be using 6-fold cross-validation, with the files being titled `SettingC/CVSpplits/training_0X.data` where X is a number between 0 and 5 (inclusive). These files, along with `SettingC/training.data` and `SettingC/test.data` can be found in the `SettingC` folder.

1. [5 points] Update your decision tree implementation to have functionality to deal with missing features. Describe your approach/any choices you had to make in this implementation.
2. [10 points] Perform 6-fold cross-validation on each of the 3 methods described above. Report the accuracy for each method and the standard deviation.
3. [5 points] Using the best method selected from your experiments, train your decision tree on `SettingC/training.data`, and report the accuracy of your tree on `SettingC/test.data`.

References

- [1] Mitchell, Tom M. “Decision Tree Learning.” *Machine Learning*. New York: McGraw-Hill, 1997. N. pag. Print.